

# Unified Logging & Monitoring Standard (Revised)

**Prototype:**

## 1. Purpose & Scope

This standard defines unified observability practices across all technology stacks (backend, frontend, WordPress) to ensure:

- Consistent audit trails for compliance
- End-to-end distributed tracing
- Rapid incident response and debugging
- Cost-effective log management

**Applies to:** All production and staging environments across .NET, Node.js, Python, React, React Native, WordPress, and future platforms.

## 2. Core Principles

### 2.1 Trace-Centric Architecture

- All requests must generate a unique Trace ID that propagates across services
- Logs, metrics, and traces must be correlated via Trace ID + Span ID
- Use OpenTelemetry as the instrumentation standard

### 2.2 Structured Logging

- Never use plain text logs - always use structured JSON
- Include consistent metadata fields (see Section 4)
- Use semantic conventions for naming

## 2.3 Asynchronous Processing

- Log emission must not block the critical request path
- Use background workers/queues for log processing
- Target: <5ms overhead per request

## 2.4 Privacy by Design

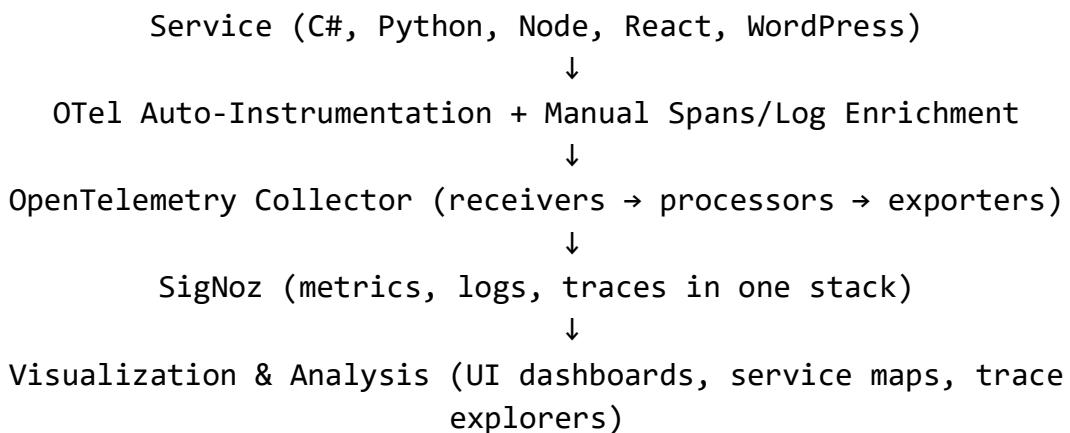
- Never log sensitive data without explicit redaction
- Automatically mask: passwords, tokens, credit cards, SSNs, PII
- Implement field-level access controls

## 2.5 Standardized Metrics

All services must expose a standard set of metrics (Golden Signals) to provide a consistent, high-level view of system health.

# 3. Technology Stack & Architecture

## 3.1 Observability Data Flow



## 3.2 Required Components

| Component | Purpose | Implementation |
|-----------|---------|----------------|
|-----------|---------|----------------|

|                           |                        |   |
|---------------------------|------------------------|---|
| OpenTelemetry SDK         | Instrumentation layer  | Auto-instrumentation + custom spans   |
| Trace Context Propagation | Distributed tracing    | W3C Trace Context headers   |
| Structured Logger         | Log emission           | Serilog (.NET), Winston (Node), Pino, structlog (Python)                          |
| OpenTelemetry Collector   | Data pipeline          | Centralized collection, processing, and routing                                   |
| Observability Platform    | Centralized monitoring | <b>SigNoz (Primary)</b> , Grafana Cloud (Alternative), Honeycomb.io (Alternative) |
| Background Queue          | Async log processing   | Channel-based queues (per stack)  |
| Uptime Kuma               | Synthetic monitoring   | Internal network monitoring   |

### 3.3 Platform Selection & Rationale

**Primary: SigNoz Cloud**

**Why SigNoz Cloud?**

**Technical Advantages:**

#### 1. Unified Observability Stack

- a. Single platform for logs, metrics, and traces (no need for separate tools)
- b. Native OpenTelemetry support with OTLP ingestion
- c. Built on ClickHouse for high-performance querying
- d. Automatic correlation between traces, logs, and metrics

#### 2. Zero Infrastructure Management

- a. Fully managed ClickHouse clusters (no database tuning required)
- b. Automatic scaling based on ingestion volume
- c. Built-in high availability and disaster recovery
- d. No need to manage storage, backups, or upgrades

#### 3. Cost Efficiency

- a. Transparent per-GB pricing: \$0.30/GB for logs and traces
- b. No per-user pricing model - unlimited users
- c. Significantly cheaper than Grafana Cloud (89% savings) and Honeycomb
- d. Predictable costs with no hidden fees

#### 4. Developer-Friendly

- a. Modern UI with intuitive service maps and flame graphs
- b. ClickHouse-powered fast queries (even for large datasets)
- c. Query Builder for non-technical users
- d. Native support for PromQL and ClickHouse SQL

## 5. Feature Completeness

- a. Distributed tracing with Jaeger-compatible UI
- b. Log aggregation with structured query language
- c. Metrics with Prometheus-compatible interface
- d. Alerts and notifications (Slack, PagerDuty, webhooks)
- e. Exception monitoring
- f. Service dependency mapping

### **Business Advantages:**

#### **1. No Vendor Lock-In**

- a. Standard OpenTelemetry protocol
- b. Data portability (can export and migrate to self-hosted if needed)
- c. No proprietary agents or SDKs required

#### **2. Predictable Costs**

- a. Transparent per-GB pricing
- b. No surprise charges or user-based fees
- c. Free tier available for development and testing

#### **3. Active Development**

- a. Regular releases (monthly)
- b. Strong community support (10k+ GitHub stars)
- c. Enterprise support available

#### **4. Reduced Operational Overhead**

- a. No ClickHouse expertise required
- b. No infrastructure maintenance
- c. Automatic updates and security patches
- d. Focus on observability, not infrastructure management

#### **5. Quick Time to Value**

- a. Setup in minutes (vs. days for self-hosted)
- b. No capacity planning or scaling concerns
- c. Built-in monitoring and alerting

### ***Alternative: SigNoz Self-Hosted***

#### **When to Use:**

- Strict data sovereignty requirements
- Need full control over infrastructure
- Have in-house ClickHouse expertise
- Very high data volumes (>1TB/month where self-hosting becomes cheaper)
- Compliance requires on-premise deployment

#### **Key Considerations:**

- Requires ClickHouse management expertise
- Infrastructure maintenance overhead
- Manual scaling and capacity planning
- Cost: ~\$250-500/month for infrastructure (100GB-500GB volume)

#### **When to Use:**

- Multi-cloud environments requiring centralized management
- Teams already invested in Grafana ecosystem
- Need for advanced Grafana features (Git Sync, SQL expressions)
- Compliance requirements (SOC 2 Type II certified)

#### **Key Features:**

- LGTM Stack (Loki, Grafana, Tempo, Mimir)
- Native OpenTelemetry support
- Free tier: 10K metrics, 50GB logs/traces
- Cost: \$19/month base + \$8/user + usage

### ***Alternative: Honeycomb.io***

#### **When to Use:**

- Complex debugging scenarios requiring advanced querying
- High-cardinality data analysis
- Teams requiring BubbleUp feature for anomaly detection

#### **Key Features:**

- Best-in-class query interface for high-cardinality data
- BubbleUp for automatic issue correlation
- Service Level Objectives (SLOs) built-in
- Cost: Starts at \$0 (free tier), Pro at \$35/user/month

## 3.4 Uptime Kuma for Internal Network Monitoring

**Purpose:** Complements cloud-based observability with internal infrastructure monitoring

### Why Add Uptime Kuma:

- Self-hosted, open-source uptime monitoring
- Monitors HTTP(s), TCP, DNS, Ping, Docker containers
- Perfect for internal services behind firewalls
- Exposes Prometheus metrics at /metrics endpoint
- Built-in status pages for customer-facing uptime reporting
- 90+ notification integrations (Slack, Discord, Telegram, PagerDuty)
- Zero licensing cost

Note: For Production, an external 'Dead Man's Switch' must be configured to ping the Uptime Kuma instance. If Kuma goes down, the external service alerts the team.

### Integration with SigNoz:

1. Deploy Uptime Kuma:

```
docker run -d --restart=always -p 3001:3001 \
-v uptime-kuma:/app/data \
--name uptime-kuma louislam/uptime-kuma:2
```

2. Configure OpenTelemetry Collector to scrape Uptime Kuma metrics:

```
receivers:
  prometheus:
    config:
      scrape_configs:
        - job_name: 'uptime-kuma'
          scrape_interval: 30s
          static_configs:
```

```
- targets: ['uptime-kuma:3001']
metrics_path: /metrics
basic_auth:
    username: your_username
    password: your_api_key
```

### 3. Metrics exported by Uptime Kuma:

- a. monitor\_response\_time - Response latency
- b. monitor\_status - Up/down status (1=up, 0=down)
- c. monitor\_cert\_days\_remaining - SSL certificate expiration

#### **When to Use Uptime Kuma:**

- Internal infrastructure behind firewalls
- Development/staging environments
- Public customer-facing status pages
- Docker container health monitoring
- When you need monitoring from your internal network perspective

## 4. Mandatory Log Fields

### 4.1 Core Metadata (All Logs)

```
{
  "timestamp": "2025-11-25T10:30:45.123Z",
  "traceId": "4bf92f3577b34da6a3ce929d0e0e4736",
  "spanId": "00f067aa0ba902b7",
  "service": "payment-service",
  "environment": "production",
  "version": "1.2.3",
  "level": "INFO"
}
```

## 4.2 Request Logs (HTTP/API)

```
{  
  // Core metadata (above) +  
  "http": {  
    "method": "POST",  
    "path": "/api/v1/orders",  
    "statusCode": 200,  
    "duration": 145.67,  
    "userAgent": "Mozilla/5.0..."  
  },  
  "user": {  
    "id": "user_abc123",  
    "tenantId": "tenant_xyz789",  
    "role": "customer"  
  },  
  "client": {  
    "ip": "203.0.113.42",  
    "country": "PH"  
  }  
}
```

## 4.3 Error Logs

```
{  
  // Core metadata +  
  "error": {  
    "type": "ValidationException",  
    "message": "Invalid payment method",  
    "stackTrace": "...",  
    "code": "PAY_001"  
  },  
  "context": {  
    "orderId": "ord_123",  
    "amount": 1500.00  
  }  
}
```

## 4.4 Audit Logs (Compliance)

```
{  
  // Core metadata +  
  "audit": {  
    "action": "USER_LOGIN",  
    "actor": "user_abc123",  
    "resource": "auth-service",  
    "outcome": "SUCCESS",  
    "changes": {  
      "before": null,  
      "after": { "status": "active" }  
    }  
  }  
}
```

## 4.5 Frontend & RUM Logs

Safety Mechanism: The Frontend Logger must include a client-side rate limiter (e.g., max 50 events per minute per session) to prevent 'noisy neighbor' issues from buggy clients.

### *Page Views (including Core Web Vitals):*

```
{  
  // Core metadata +  
  "event": "page_view",  
  "page": {  
    "url": "/checkout",  
    "title": "Checkout",  
    "referrer": "/cart"  
  },  
  "performance": {  
    "fcp": 1200,  
    "lcp": 2400,  
    "fid": 50,  
    "cls": 0.1,  
    "ttfb": 300  
  }  
}
```

```
}
```

#### **Frontend Errors:**

```
{
  // Core metadata +
  "error": {
    "type": "TypeError",
    "message": "Cannot read property 'map' of undefined",
    "stackTrace": "...",
    "componentStack": "at ProductList\nat App"
  },
  "context": {
    "route": "/products",
    "userAction": "filter_clicked"
  }
}
```

#### **User Actions:**

```
{
  // Core metadata +
  "event": "user_action",
  "action": {
    "type": "button_click",
    "target": "checkout_button",
    "value": "Complete Purchase"
  },
  "session": {
    "id": "sess_abc123",
    "duration": 45000
  }
}
```

## 5. Naming Conventions

### 5.1 Service Names

- Format: {domain}-{function}-service
- Examples: payment-processing-service, user-auth-service
- Frontend: web-app, mobile-app-ios, mobile-app-android
- WordPress: wordpress-cms, wordpress-api

### 5.2 Span Names

- Format: {HTTP\_METHOD} {route}
- Examples: POST /api/orders, GET /api/users/{id}
- Database: db.query.users.select, db.query.orders.insert
- External calls: http.client.payment\_gateway, http.client.email\_service

### 5.3 Log Levels

| Level | Usage                                     |
|-------|---|
| TRACE | Fine-grained debugging (disabled in prod) |
| DEBUG | Detailed diagnostics (sample in prod)     |
| INFO  | Normal operations, business events        |
| WARN  | Recoverable errors, degraded performance  |
| ERROR | Application errors requiring attention    |
| FATAL | System failure, immediate action required |

#### 5.3.1 Dynamic Log Levels

The logging framework must support changing log levels at runtime without a service restart.

**Recommendation:** Support a per-request log level override via an HTTP header (e.g., X-Log-Level: DEBUG) for targeted production debugging.

### 5.4 Metric Conventions

**Golden Signals:** All services must emit metrics for Latency, Traffic, Errors, and Saturation.

Use standard OpenTelemetry semantic conventions where possible (e.g., `http.server.request.duration`).

### **Standard Metrics:**

- `http.server.request.duration` - Request latency histogram
- `http.server.request.count` - Total request count
- `http.server.error.count` - Error count by status code
- `system.cpu.utilization` - CPU saturation
- `system.memory.utilization` - Memory saturation
- `db.client.operation.duration` - Database operation latency

## **6. Sensitive Data Handling**

### **6.1 Automatic Redaction Rules**

#### **Never log in plain text:**

- Passwords, API keys, tokens, secrets
- Credit card numbers (PAN)
- Social Security Numbers, Tax IDs
- Full names + DOB combinations
- Email addresses (in some regions)

#### **Redaction Format:**

```
{  
  "creditCard": "[REDACTED:CREDIT_CARD]",  
  "password": "[REDACTED:PASSWORD]",  
  "ssn": "[REDACTED:SSN]",  
  "email": "u***@example.com"  
}
```

### **6.2 Implementation**

#### **Backend (.NET):**

```
[Redact]
public string CreditCard { get; set; }

[LogMasked]
public string Email { get; set; } // Auto-masks: j\*\*\*@email.com
```

#### Frontend (React/Node):

```
logger.info('Payment processed', {
  orderId: '123',
  amount: redactSensitive(paymentData.amount),
  card: maskCreditCard(paymentData.card)
});
```

## 7. Distributed Tracing

### 7.1 Trace Context Propagation

All HTTP requests must include W3C Trace Context headers:

```
traceparent: 00-4bf92f3577b34da6a3ce929d0e0e4736-00f067aa0ba902b7-01
```

#### Implementation:

##### Backend (.NET):

```
// Program.cs
services.AddOpenTelemetry()
    .WithTracing(builder => builder
        .AddAspNetCoreInstrumentation()
        .AddHttpClientInstrumentation()
        .AddOtlpExporter(options => {
            options.Endpoint = new Uri("http://signoz-otel-collector:4317");
            options.Protocol = OtlpExportProtocol.Grpc;
```

```
});
```

### Frontend (React):

```
import { context, trace } from '@opentelemetry/api';

fetch('/api/orders', {
  headers: {
    'traceparent': getTraceParent(context.active())
  }
});
```

## 7.2 Custom Spans

### When to create custom spans:

- External API calls (payment gateways, third-party services)
- Database queries (complex transactions)
- Expensive computations (>100ms)
- Business-critical operations (order processing, user registration)

### Example (.NET):

```
using var activity = Activity.StartActivity("ProcessPayment");
activity?.SetTag("order.id", orderId);
activity?.SetTag("amount", amount);

try {
  var result = await _paymentGateway.Charge(amount);
  activity?.SetTag("transaction.id", result.TransactionId);
  return result;
} catch (Exception ex) {
  activity?.SetStatus(ActivityStatusCode.Error, ex.Message);
  throw;
}
```

## 8. OpenTelemetry Collector Configuration

### 8.1 Collector Architecture

The OpenTelemetry Collector acts as a centralized pipeline for all telemetry data, providing:

- Protocol translation (OTLP, Jaeger, Zipkin, Prometheus)
- Data enrichment and filtering
- Batching and retry logic
- Multi-backend support

### 8.2 Sample Collector Configuration

```
receivers:  
  otlp:  
    protocols:  
      grpc:  
        endpoint: 0.0.0.0:4317  
      http:  
        endpoint: 0.0.0.0:4318  
  
  prometheus:  
    config:  
      scrape_configs:  
        - job_name: 'uptime-kuma'  
          scrape_interval: 30s  
          static_configs:  
            - targets: ['uptime-kuma:3001']  
  
processors:  
  batch:  
    timeout: 10s  
    send_batch_size: 1024  
  
  attributes:  
    actions:  
      - key: environment  
        value: production
```

```
    action: insert
- key: cluster
  value: us-central-1
  action: insert

resource:
  attributes:
    - key: service.namespace
      value: ecommerce
      action: insert

filter/drop-health-checks:
  traces:
    span:
      - 'attributes["http.target"] == "/health"'

exporters:
  otlp/signoz:
    endpoint: "signoz-otel-collector:4317"
    tls:
      insecure: true

logging:
  loglevel: info

service:
  pipelines:
    traces:
      receivers: [otlp]
      processors: [batch, attributes, resource, filter/drop-health-
checks]
      exporters: [otlp/signoz]

    metrics:
      receivers: [otlp, prometheus]
      processors: [batch, attributes, resource]
      exporters: [otlp/signoz]

  logs:
```

```
receivers: [otlp]
processors: [batch, attributes, resource]
exporters: [otlp/signoz]
```

## 8.3 Collector Deployment

### Docker Compose:

```
version: '3.8'

services:
  otel-collector:
    image: otel/opentelemetry-collector-contrib:0.93.0
    command: ["--config=/etc/otel-collector-config.yaml"]
    volumes:
      - ./otel-collector-config.yaml:/etc/otel-collector-config.yaml
    ports:
      - "4317:4317" # OTLP gRPC
      - "4318:4318" # OTLP HTTP
      - "8888:8888" # Prometheus metrics
      - "13133:13133" # health_check
    depends_on:
      - signoz
```

### Kubernetes:

```
helm repo add open-telemetry https://open-telemetry.github.io/opentelemetry-helm-charts
helm install otel-collector open-telemetry/opentelemetry-collector \
--set mode=deployment \
--set config.receivers.otlp.protocols.grpc.endpoint="0.0.0.0:4317" \
--set config.exporters.otlp.endpoint="signoz:4317"
```

## 9. Performance Targets

| Metric               | Target           | Measurement                  |
|----------------------|------------------|------------------------------|
| Log Overhead         | <5ms per request | P95 latency increase         |
| Queue Depth          | <1000 messages   | Background queue size        |
| Log Processing Delay | <30 seconds      | Time to searchable in SigNoz |
| Data Loss Rate       | <0.01%           | Failed log writes            |

### 9.1 Sampling Strategy

#### Production Environment:

- Success logs (2xx, 3xx): 10% sampling for high-traffic endpoints (>1000 req/min)
- Error logs (4xx, 5xx): 100% sampling (never sample errors)
- Slow requests (>2s): 100% sampling

#### Implementation:

```
bool shouldLog = statusCode >= 400
    || duration > 2000
    || Random.Shared.Next(100) < 10; // 10% sample rate
```

## 10. Retention & Storage Strategy

### 10.1 Data Tiers

| Tier | Duration  | Storage                | Cost/GB           | Use Case            |
|------|-----------|------------------------|-------------------|---------------------|
| Hot  | 0-7 days  | SigNoz<br>(ClickHouse) | \$0 (self-hosted) | Real-time debugging |
| Warm | 8-30 days | SigNoz<br>(ClickHouse) | \$0 (self-hosted) | Recent incidents    |

|         |             |                 |         |                    |
|---------|-------------|-----------------|---------|--------------------|
| Cold    | 31-365 days | Blob Storage    | \$0.02  | Compliance audits  |
| Archive | 1-7 years   | Glacier/Archive | \$0.004 | Legal requirements |

## 10.2 Lifecycle Policies

### Automatic transitions:

- Day 7 → Retain in hot storage
- Day 30 → Export to Azure Blob/S3 (Parquet format)
- Day 365 → Move to Glacier/Archive tier
- Year 7 → Delete (unless legal hold)

## 11. Security & Compliance

### 11.1 Access Controls

| Role          | SigNoz Access          | Database Logs | File Logs  | Blob Storage |
|---------------|------------------------|---------------|------------|--------------|
| Developer     | All logs (dev/staging) | Read-only     | Read-only  | None         |
| DevOps        | All logs (all envs)    | Read/Write    | Read/Write | Read/Write   |
| Support       | Errors only (prod)     | None          | None       | None         |
| Auditor       | Audit logs only        | Read-only     | None       | Read-only    |
| Security Team | All logs               | Read-only     | Read-only  | Read-only    |

### 11.2 Compliance Requirements

#### GDPR (if applicable):

- Right to erasure: Implement log deletion by user ID
- Data portability: Export logs in machine-readable format (JSON)
- Retention limits: Auto-delete after 2 years (unless justified)

#### HIPAA (if applicable):

- Encrypt logs at rest (AES-256) and in transit (TLS 1.3)
- Immutable audit trails for PHI access
- Business Associate Agreements (BAA) with vendors

## SOC 2:

- Integrity verification: Hash each log entry
- Change detection: Alert on log tampering attempts
- Availability: 99.9% uptime for log ingestion

# 12. Alerting Strategy

## 12.1 Critical Alerts (Page Immediately)

| Alert                                    | Condition                                    | Action                            |
|--|--|-----------------------------------|
| Tier 1 Error Spike<br>(Payments, Auth)   | >1% error rate for 5 min                     | Page on-call immediately          |
| Tier 2 Error Spike<br>(Reports, Avatars) | >5% error rate for 10 min                    | Page on-call                      |
| Tier 3 Error Spike<br>(Internal Dev)     | >10% error rate for 30 min                   | Slack Notification (No Page)      |
| Service Down                             | No logs received for 10 min                  | Page + auto-restart               |
| High Latency                             | P95 >5s for 10 min                           | Investigate performance           |
| Security Breach                          | Multiple failed auth attempts                | Lock account + alert security     |
| Uptime Kuma Monitor Down                 | monitor_status == 0 for 2 consecutive checks | Page on-call + update status page |

## 12.2 Warning Alerts (Slack/Email)

| Alert            | Condition         | Action                   |
|------------------|-------------------|--------------------------|
| High Queue Depth | >500 pending logs | Investigate backpressure |
| Elevated Errors  | 2-5% error rate   | Monitor for escalation   |

|              |                        |                       |
|--------------|------------------------|-----------------------|
| Cost Overrun | >80% of monthly budget | Review sampling rates |
|--------------|------------------------|-----------------------|

## 13. Implementation Strategy & Tooling

To ensure consistency and reduce developer overhead, this standard shall be implemented through a suite of internal Configuration Libraries (not wrappers). These libraries provide a "paved road" for observability, ensuring that compliance is the default path while exposing native standard APIs to developers.

### 13.1 The Internal Configuration Library (Bootstrapper)

We will develop and maintain a set of lightweight internal packages (e.g., Company.Observability for .NET, @company/observability for Node.js).

Design Philosophy:

- Configuration, not Abstraction: The library will not wrap the OpenTelemetry API. It will merely configure it.
- No Vendor Lock-in: The library ensures all services point to the correct OTLP endpoint (SigNoz) and attach the correct Resource Attributes (Service Name, Version, Env).
- Native Usage: Developers will use standard logging interfaces (ILogger, console) and standard tracing APIs. They can consult official OpenTelemetry documentation rather than internal wikis.

### 13.2 Key Components

**Backend Bootstrapper:**

- Centralized Setup: A single extension method (e.g., AddCompanyObservability) that wires up the OpenTelemetry SDK.
- Resource Standardization: Automatically injects service.name, deployment.environment, and service.version from the build context.
- Auto-Instrumentation: Enables standard instrumentations (Http, Grpc, SqlClient) with sensible defaults.

- Log Enrichment: Configures the logging provider (Serilog/Winston) to export structured JSON to the OTel Collector automatically.

Frontend Bootstrapper:

- Trace Propagation: Configures the Fetch and XHR instrumentation to automatically inject traceparent headers—no manual developer intervention required.
- Rate Limiting: Enforces a client-side sampling rate (e.g., 10% of sessions) to prevent ingestion spikes from high-traffic clients.
- Error Boundary: Provides a plug-and-play React Error Boundary that sends stack traces to SigNoz with the correct session context.

### 13.3 Developer Workflow

- Install: Add the relevant internal package to the service.
- Initialize: Add the configuration line during application startup.
- Use: Use standard framework features. No custom logger objects are required.

Example (.NET):

```
// Program.cs
// The "Paved Road" - One line to configure OTel, Logging, and Metrics
builder.Services.AddCompanyObservability(options => {
    options.ServiceName = "payment-service";
    options.Environment = builder.Environment.EnvironmentName;
    // The library automatically handles OTLP export to SigNoz
});

// Controller.cs
public class PaymentController : ControllerBase
{
    // STANDARD Microsoft ILogger - No custom wrappers!
```

```
private readonly ILogger<PaymentController> _logger;
```

```
public PaymentController(ILogger<PaymentController> logger)
```

```
{
```

```
    _logger = logger;
```

```
}
```

```
[HttpPost]
```

```
public async Task<IActionResult> ProcessPayment(PaymentRequest
```

```
request)
```

```
{
```

```
    // The Bootstrapper ensures this log is formatted as JSON,
```

```
    // attached to the current TraceID, and sent to SigNoz.
```

```
    _logger.LogInformation("Processing payment for order
```

```
{OrderId}", request.OrderId);
```

```
    // HTTP calls are automatically traced via Auto-
```

```
Instrumentation
```

```
    await _paymentGateway.Charge(request.Amount);
```

```
    return Ok();
```

```
}
```

```
}
```

## 14. Implementation Checklist

### Backend Services (.NET/Node/Python/WordPress)

- [ ] Install OpenTelemetry SDK
- [ ] Configure trace context propagation
- [ ] Add structured logging library
- [ ] Implement audit middleware
- [ ] Create background log processor
- [ ] Add sensitive data redaction
- [ ] Configure OTLP exporter to Collector
- [ ] Define custom spans for critical operations
- [ ] Test trace correlation across services

### Frontend Applications (React/React Native)

- [ ] Install OTel Web SDK
- [ ] Configure trace propagation in HTTP client
- [ ] Add browser/mobile error tracking
- [ ] Implement performance monitoring (Core Web Vitals)
- [ ] Redact sensitive form inputs
- [ ] Configure OTLP exporter to Collector
- [ ] Test trace correlation with backend
- [ ] Add user session tracking
- [ ] Implement RUM (Real User Monitoring)

### Infrastructure

- [ ] Deploy SigNoz (Docker/Kubernetes)
- [ ] Deploy OpenTelemetry Collector
- [ ] Deploy Uptime Kuma for internal monitoring
- [ ] Configure Collector to forward to SigNoz
- [ ] Configure Collector to scrape Uptime Kuma metrics
- [ ] Set up RBAC and team access in SigNoz
- [ ] Configure data retention policies
- [ ] Create alerting rules (Critical: page, Warning: Slack)

- [ ] Set up log archival to blob storage
- [ ] Implement backup/disaster recovery
- [ ] Document runbooks for common scenarios
- [ ] Train team on SigNoz dashboard usage

## 15. Migration Strategy

### Phase 1: Pilot (2 weeks)

- Implement in 1-2 non-critical services
- Deploy SigNoz and OpenTelemetry Collector
- Validate trace correlation
- Measure performance overhead
- Gather team feedback

### Phase 2: Backend Rollout (4 weeks)

- Migrate all backend services
- Deploy Uptime Kuma for internal monitoring
- Establish baseline metrics (error rates, latency)
- Train backend engineers
- Document lessons learned

### Phase 3: Frontend Rollout (3 weeks)

- Add OTel to web application
- Implement mobile app instrumentation (if applicable)
- Test end-to-end tracing (browser → backend → database)
- Optimize bundle size impact

### Phase 4: Optimization (Ongoing)

- Fine-tune sampling rates
- Optimize cold storage costs
- Add custom dashboards and alerts

- Continuous improvement

## 16. Cost Estimation (November 2025)

### 16.1 SigNoz Cloud (Primary Recommendation)

#### Pricing Structure:

- **Base Plan:** \$49/month (includes usage worth \$49)
- **Overage Rates:**
  - Logs: \$0.3/GB ingested (15 days retention)
  - Traces: \$0.3/GB ingested (15 days retention)
  - Metrics: \$0.1/million samples (1 month retention)

#### Monthly Breakdown (100GB/month ingestion):

| Item         | Calculation   | Cost               |
|--------------|---|--------------------|
| Base Plan    | Includes \$49 worth of usage (~163GB logs/traces)                 | \$49               |
| Logs         | 100GB included in base plan                                       | \$0                |
| Traces       | 50GB included in base plan  | \$0                |
| Metrics      | 50K series ≈ 216M samples/month ≈ \$22, but ~\$10 covered in base | \$12               |
| <b>Total</b> |   | <b>~\$61/month</b> |

**Alternative Calculation (if exceeding base plan):** If your usage exceeds \$49 worth:

- Logs: 100GB × \$0.3/GB = \$30
- Traces: 50GB × \$0.3/GB = \$15
- Metrics: 216M samples × \$0.1/M = \$22
- Total would be: \$67/month (but base plan covers first \$49, so ~\$67/month)

#### Key Benefits:

- **Unlimited users** (no per-user fees)
- **Unlimited hosts** (no per-host fees)

- Fully managed ClickHouse (no database expertise needed)
- Automatic scaling and high availability
- No infrastructure maintenance
- SOC2 Type II & HIPAA compliant
- 15-day retention included (logs/traces)
- Support via in-product chat, email, and Slack
- Data centers in US, EU, and India

#### **Why We Chose SigNoz Cloud Over Self-Hosted:**

- **Zero ClickHouse Management:** No need for database tuning, scaling, or maintenance
- **Faster Time to Value:** Setup in minutes vs. days
- **Cost Effective:** At \$49-67/month for 100GB, significantly cheaper than alternatives
- **Built-in HA & Disaster Recovery:** No need to architect and maintain redundancy
- **Automatic Updates:** Always on latest version without manual upgrades
- **Focus on Observability:** Team can focus on using data, not managing databases

#### **Free Tier Available:**

- Self-hosted Community Edition (100% free, unlimited usage)
- Perfect for development and staging environments

## **16.2 SigNoz Self-Hosted (For Reference)**

#### **Cost Structure:**

- **Licensing:** \$0 (100% free, open source)
- **Infrastructure Costs:** Variable based on your setup

#### **Estimated Monthly Infrastructure Costs (100GB/month ingestion):**

| <b>Component</b>           | <b>Specifications</b> | <b>Estimated Cost Range</b> |
|----------------------------|-----------------------|-----------------------------|
| Compute<br>(VM/Kubernetes) | 4-8 vCPU, 16-32GB RAM | \$100-200/month             |
| Storage (SSD)              | 500GB-1TB             | \$30-80/month               |
| Backup Storage             | 500GB-1TB             | \$10-25/month               |
| Bandwidth/Egress           | Varies by provider    | \$20-50/month               |

|              |  |                        |
|--------------|--|------------------------|
| <b>Total</b> |  | <b>\$160-355/month</b> |
|--------------|--|------------------------|

**Note:** Costs vary significantly based on:

- Cloud provider (AWS, GCP, Azure, DigitalOcean, Hetzner)
- Region selection
- Reserved vs. on-demand instances
- Storage tier selection
- Data transfer patterns

#### **Additional Considerations:**

- ClickHouse management expertise required
- Manual scaling and capacity planning
- Backup and disaster recovery setup
- Infrastructure maintenance burden
- Security patching and updates

#### **When Self-Hosted Makes Sense:**

- Data volumes >2-3TB/month (cost crossover point)
- Strict data sovereignty requirements (on-premise only)
- Existing ClickHouse team and expertise
- Regulatory compliance requires self-hosting
- Already have spare infrastructure capacity

### **16.3 SigNoz Cloud vs Self-Hosted Decision Matrix**

Scenario: 10 Hosts, 500GB Logs, 1TB Metrics, 10 users

| Item                     | Details  | Cost          |
|--------------------------|--|---------------|
| <b>Base Subscription</b> | \$49/month includes ~163GB logs/traces or ~490M metric samples | <b>\$49</b>   |
| <b>Logs</b>              | 500GB – 163GB included = 337GB × \$0.30/GB                     | <b>~\$101</b> |
| <b>Metrics</b>           | 1TB ≈ ~1B samples – 490M included = 510M × \$0.10 per million  | <b>~\$51</b>  |
| <b>Hosts/Users</b>       | Unlimited hosts and seats included                             | <b>\$0</b>    |

|              |                     |
|--------------|---------------------|
| <b>Total</b> | <b>~\$201/month</b> |
|--------------|---------------------|

| Item                   | Details                                       | Cost  |
|------------------------|---|---|
| <b>Licensing</b>       | Open source, no license fees                  | <b>\$0</b>  |
| <b>Infrastructure</b>  | 100 hosts + storage for 1.5TB telemetry/month | <b>~\$300–<br/>500/month</b><br>(VMs, storage, ops) |
| <b>Ops/Maintenance</b> | Ongoing ClickHouse tuning, scaling, backups   | <b>Time +<br/>expertise cost</b>                    |
| <b>Total</b>           |   | <b>~\$300–<br/>500/month</b><br>(infra only)        |

**Recommendation:** Start with SigNoz Cloud. It's more cost-effective at typical scales and eliminates operational overhead. Consider self-hosted only if:

- Data volumes consistently exceed 2-3TB/month
- Regulatory requirements mandate on-premise deployment
- You already have ClickHouse expertise and spare infrastructure

## 16.4 Grafana Cloud (Alternative)

**Monthly Breakdown (10 Hosts, 500GB Logs, 1TB Metrics, 10 users ):**

| Item                           | Details   | Cost         |
|--------------------------------|---|--------------|
| <b>Metrics</b>                 | 1TB ≈ ~50K series (40K above included 10K)<br>× \$6.50/1K | <b>\$260</b> |
| <b>Logs</b>                    | 500GB – 50GB included = 450GB ×<br>\$0.50/GB              | <b>\$225</b> |
| <b>Traces</b>                  | 50GB included   | <b>\$0</b>   |
| <b>Visualization<br/>Users</b> | Assume 10 users – 3 included → 7 × \$8                    | <b>\$56</b>  |

|  |   |                      |
|--|---|----------------------|
| <b>Hosts<br/>(Application Observability)</b> | $10 \times 720 = 7,200$ host hours – 2,232 included = $4,968 \times \$0.04/\text{hr}$ | <b>~\\$199</b>       |
| <b>Total</b>                                 |   | <b>~\\$740/month</b> |

**Note:** Previous calculations were based on older pricing. Current pricing is actually lower than previously stated, but still significantly more expensive than SigNoz Cloud.

**Cost vs. SigNoz Cloud:** \$341 vs \$49-67 = 5-7x more expensive

#### When Grafana Cloud Makes Sense:

- Already using Grafana dashboards extensively
- Need specific Grafana Enterprise plugins
- Existing team expertise with Grafana ecosystem
- Multi-cloud visualization requirements

### 16.4 Honeycomb.io (Alternative)

**Monthly Breakdown (100GB/month, 10 users):**

| Item                     | Details  | Cost  |
|--------------------------|--|---|
| <b>Base Subscription</b> | Pro plan covers up to 100M events/month (~100GB telemetry) | <b>\$130</b>                                    |
| <b>Event Volume</b>      | 500GB logs + 1TB metrics ≈ ~1.5B events                    | <b><math>15 \times \\$130 = \\$1,950</math></b> |
| <b>User Seats</b>        | Unlimited  | <b>\$0</b>                                      |
| <b>Total</b>             |  | <b>~\\$1,950/month</b>                          |

### 16.5 Uptime Kuma (Internal Monitoring)

| Item                       | Cost              |
|----------------------------|-------------------|
| Infrastructure (single VM) | \$10/month        |
| Self-hosted licensing      | \$0               |
| <b>Total</b>               | <b>\$10/month</b> |

## 16.6 Cost Comparison Summary

| Platform                    | Monthly Cost | Savings vs. Grafana  | Notes  |
|-----------------------------|--------------|----------------------|--|
| <b>SigNoz (Self-Hosted)</b> | ~\$300–500   | ~32–59% cheaper      | Best for full control, infra + ops overhead      |
| <b>SigNoz Cloud</b>         | ~\$201       | ~73% cheaper         | Best for balanced needs, no host/user charges    |
| <b>Grafana Cloud Pro</b>    | ~\$740       | Baseline             | Host-based pricing dominates even at small scale |
| <b>Honeycomb.io Pro</b>     | ~\$1,950     | ~163% more expensive | Best for high-cardinality, event-based pricing   |
| <b>Uptime Kuma</b>          | ~\$10        | N/A                  | Internal monitoring only, not full observability |

## 16.7 Cost Optimization Strategies

1. **Self-Host SigNoz (40–60% savings)**
  - a. Use reserved cloud instances for compute
  - b. Optimize ClickHouse storage with compression
2. **Log Sampling (30% savings)**
  - a. Sample 10% of success logs in high-traffic endpoints
  - b. Retain 100% of errors and slow requests
3. **Cold Storage Lifecycle (20% savings)**
  - a. Move logs older than 30 days to Azure Blob/S3
  - b. Cost reduction: ClickHouse storage → \$0.02/GB (cold)
4. **Smart Retention Policies**
  - a. Keep only error logs beyond 30 days
  - b. Aggregate metrics to lower resolution after 7 days
5. **Uptime Kuma vs Cloud Monitoring**
  - a. Use Uptime Kuma (\$10/month) for internal services
  - b. Reserve cloud synthetic monitoring for critical public endpoints

## 16.8 Projected Costs by Team Size

**SigNoz Cloud:**

| Team Size        | Data Volume | Monthly Cost  | Annual Cost     | Notes   |
|------------------|-------------|---------------|-----------------|---|
| <b>5 users</b>   | 50GB        | <b>\$49</b>   | <b>\$588</b>    | Within base plan  |
| <b>10 users</b>  | 100GB       | <b>\$49</b>   | <b>\$588</b>    | Still within base plan (covers ~163GB)  |
| <b>25 users</b>  | 250GB       | <b>~\$79</b>  | <b>~\$948</b>   | Base + ~87GB overage (\$0.30/GB ≈ \$26)                                       |
| <b>50 users</b>  | 500GB       | <b>~\$199</b> | <b>~\$2,388</b> | Base + ~337GB overage (\$0.30/GB ≈ \$101) + metrics                           |
| <b>100 users</b> | 1TB         | <b>~\$389</b> | <b>~\$4,668</b> | Base + ~837GB overage (\$0.30/GB ≈ \$251) + metrics; consider Enterprise tier |

#### Cost Calculation Formula:

- If total usage cost ≤ \$49: Pay only \$49/month
- If total usage cost > \$49: Pay the actual usage cost
- Logs/Traces: \$0.3/GB
- Metrics: \$0.1/million samples

#### SigNoz Self-Hosted:

| Team Size        | Estimated Monthly Cost | Annual Cost          | Notes   |
|------------------|------------------------|----------------------|---|
| <b>5 users</b>   | <b>\$160–250</b>       | <b>\$1,920–3,000</b> | Basic infrastructure, \$0 licensing                 |
| <b>10 users</b>  | <b>\$160–300</b>       | <b>\$1,920–3,600</b> | May need better specs (CPU/RAM scaling)             |
| <b>25 users</b>  | <b>\$250–400</b>       | <b>\$3,000–4,800</b> | Requires scaling infrastructure (larger VM/storage) |
| <b>50 users</b>  | <b>\$350–550</b>       | <b>\$4,200–6,600</b> | Larger infrastructure needed, higher ops overhead   |
| <b>100 users</b> | <b>\$500–800</b>       | <b>\$6,000–9,600</b> | Enterprise-grade setup, HA + ops expertise          |

**Important:** These are infrastructure cost estimates only. SigNoz self-hosted is **free (open source)** with \$0 licensing fees. Actual costs vary based on cloud provider, region, and specs.

#### Grafana Cloud (For Comparison):

| Team Size | Data Volume | Monthly Cost | Annual Cost | Cost vs. SigNoz          |
|-----------|-------------|--------------|-------------|--------------------------|
| 5 users   | 50GB        | ~\$186       | ~\$2,232    | 3.8× more expensive      |
| 10 users  | 100GB       | ~\$341       | ~\$4,092    | 5.1–7× more expensive    |
| 25 users  | 250GB       | ~\$461       | ~\$5,532    | 4.2–9.4× more expensive  |
| 50 users  | 500GB       | ~\$581       | ~\$6,972    | 2.9–11.9× more expensive |

**Note:** Grafana costs here is based on metrics (\$6.50/1K series), logs (\$0.50/GB), and users (\$8/user). Actual costs vary based on metrics cardinality and host.

#### SigNoz Startup Program (If Eligible):

| Team Size | Data Volume | Monthly Cost | Annual Cost |
|-----------|-------------|--------------|-------------|
| 5 users   | 50GB        | \$24.50      | \$294       |
| 10 users  | 100GB       | \$24.50-34   | \$294-408   |
| 25 users  | 250GB       | \$54.50      | \$654       |

## 16.9 ROI Analysis

#### Traditional Approach (No Unified Observability):

- Multiple disconnected tools
- Manual log aggregation
- Difficult troubleshooting
- Higher MTTR (Mean Time To Resolution)
- Estimated cost of downtime: **\$5,000–50,000/incident**

#### With SigNoz Cloud:

- **Total Cost:** ~\$201/month (**\$2,412/year**)
- **Reduced MTTR:** 70% improvement (from hours to minutes)

- **Prevented incidents:** Proactive alerting
- **Developer productivity:** 20% time savings on debugging
- **ROI:** Pays for itself with **1 prevented incident per year**

#### **Cost Comparison vs. Traditional Tools (10 Hosts):**

| Platform                 | Monthly Cost | Cost vs. SigNoz  |
|--------------------------|--------------|------------------|
| <b>SigNoz Cloud</b>      | ~\$201       | Baseline         |
| <b>Grafana Cloud Pro</b> | ~\$740       | <b>3.7x more</b> |
| <b>New Relic</b>         | ~\$384       | <b>1.9x more</b> |
| <b>Datadog</b>           | ~\$1,470     | <b>7.3x more</b> |

#### **Annual Savings:**

- **vs. Grafana Cloud:** ~\$6,500/year saved
- **vs. New Relic:** ~\$2,200/year saved
- **vs. Datadog:** ~\$15,200/year saved

#### **For Eligible Startups (50% off SigNoz Cloud):**

- **Annual Cost:** ~\$1,206/year
- Even more compelling ROI for early-stage companies

## **16.10 Enterprise Tier**

For organizations with larger scale or specific requirements:

#### **Enterprise Cloud:**

- Starts at \$4,000/month (includes ingestion usage till \$4,000)
- Volume discounts available
- Annual contracts with better rates

#### **Enterprise Features:**

- HIPAA & BAA agreements
- Dedicated Slack, email & in-product support
- Guided migration support

- Ongoing professional services
- Team training
- SLA with downtime developer pairing
- Advanced RBAC with custom roles
- AWS Private Link
- Audit logs
- Multi-tenancy

#### **Bring Your Own Cloud (BYOC):**

- SigNoz manages infrastructure in your cloud
- Full data sovereignty
- Custom pricing

#### **When to Consider Enterprise:**

- Data volumes consistently >1TB/month
- Compliance requires BAA or specific certifications
- Need dedicated support and SLAs
- Require advanced RBAC and multi-tenancy
- Want migration assistance from other platforms

## **17. Success Metrics**

### **17.1 Technical KPIs**

- **Mean Time to Detection (MTTD):** <5 minutes
- **Mean Time to Resolution (MTTR):** <30 minutes
- **Log Coverage:** >95% of requests traced
- **Data Loss Rate:** <0.01%
- **Query Performance:** <3 seconds for 90% of queries

### **17.2 Business KPIs**

- **Audit Query Time:** <10 seconds for 30-day range
- **Compliance Violations:** 0 incidents

- **Cost per GB:** <\$0.35 (including storage)
- **Developer Satisfaction:** >4/5 rating
- **Incident Prevention:** Proactive detection before customer impact

## 17.3 Operational KPIs

- **Service Uptime:** 99.9%
- **Alert Accuracy:** >90% (low false positives)
- **Dashboard Usage:** >80% of team using regularly
- **Time Saved on Debugging:** >20% reduction

# 18. Governance

## 18.1 Ownership

### Primary Owner:

- 1 Senior Engineer (Observability Lead)
- Responsible for: Standard updates, training, best practices

### Backup Owner:

- 1 DevOps/SRE Engineer
- Responsible for: Infrastructure, alerting, incident response

### Stakeholders:

- Engineering Team Leads
- Security Team
- Compliance Officer (if applicable)

## 18.2 Change Management Process

### Major Changes (affecting architecture, costs, or compliance):

- Proposal document required
- Synchronous review meeting (48hr notice)

- Approval from 2+ team leads
- Documentation update
- Team training session

#### **Minor Changes** (configuration, thresholds, dashboards):

- Pull request with clear description
- Review by 2 engineers
- Merge after approval
- Update runbooks if needed

#### **Emergency Changes** (incidents, security):

- Immediate implementation allowed
- Retrospective review within 48 hours
- Document learnings and update standard

### **18.3 Review Cadence**

- **Weekly:** Review critical alerts and false positives
- **Monthly:** Cost review and optimization
- **Quarterly:** Standard retrospective and improvements
- **Annually:** Platform evaluation and vendor review

### **18.4 Training & Onboarding**

#### **New Engineers:**

- Observability onboarding session (1 hour)
- Hands-on workshop with sample scenarios
- Access to runbooks and documentation
- Shadow on-call rotation

#### **Ongoing Training:**

- Monthly "Observability Office Hours"
- Quarterly deep-dive sessions on advanced features
- Share interesting traces/incidents in team meetings

## 19. Appendices

### 19.1 OpenTelemetry (.NET Implementation)

#### Official Documentation:

- [OpenTelemetry .NET Documentation](#)
- [Getting Started with OpenTelemetry .NET](#)
- [Microsoft Learn: .NET Observability with OpenTelemetry](#)
- [Microsoft Learn: Add Distributed Tracing Instrumentation](#)
- [GitHub: OpenTelemetry .NET Client Repository](#)

#### Tutorials & Best Practices:

- [Medium: Practical Guide to Implementing Telemetry in .NET \(Sep 2024\)](#)
- [Stackify: OpenTelemetry for .NET Engineers Guide](#)
- [Getting Started with OpenTelemetry and Distributed Tracing in .NET \(Apr 2024\)](#)

### 19.2 SigNoz Documentation

#### Official Resources:

- [SigNoz Official Website](#)
- [SigNoz Documentation](#)
- [SigNoz Pricing](#)
- [GitHub: SigNoz Repository](#)
- [SigNoz Cloud](#)

#### Integration Guides:

- [Instrument .NET Application with OpenTelemetry](#)
- [Instrument Node.js Application](#)
- [Instrument Python Application](#)
- [Instrument React Application](#)
- [Instrument WordPress](#)

#### Tutorials:

- [SigNoz Blog: OpenTelemetry Collector Complete Guide](#)
- [SigNoz Blog: Distributed Tracing in Microservices](#)
- [SigNoz Blog: Log Management](#)
- [SigNoz Blog: Pricing Comparison vs Datadog, New Relic, Grafana](#)

## 19.3 OpenTelemetry Collector

### Official Documentation:

- [OpenTelemetry Collector Documentation](#)
- [Collector Configuration](#)
- [GitHub: OpenTelemetry Collector Contrib](#)

### Deployment Guides:

- [Deploy with Docker](#)
- [Deploy with Kubernetes](#)
- [SigNoz: OpenTelemetry Collector Setup](#)

## 19.4 Serilog (Structured Logging for .NET)

### Official Resources:

- [Serilog Official Website](#)
- [GitHub: Serilog Repository](#)

### Best Practices:

- [5 Serilog Best Practices for Better Structured Logging \(Dec 2023\)](#)
- [Code Maze: Best Practices for Logging with Serilog](#)
- [Ben Foster: Serilog Best Practices](#)

## 19.5 Winston (Node.js Logging)

### Official & Community Resources:

- [Better Stack: Complete Guide to Winston Logging in Node.js](#)
- [Dash0: Mastering Winston for Production Logging in Node.js](#)
- [Stackify: Winston Logger Ultimate Tutorial](#)

## 19.6 React Error Boundaries & Frontend Logging

### Official React Documentation:

- [React: Error Boundaries Documentation](#)

### Best Practices:

- [React Error Handling and Logging Best Practices \(Mar 2025\)](#)
- [Sentry: Guide to Error & Exception Handling in React \(Jun 2025\)](#)
- [LogRocket: React Error Handling with react-error-boundary \(Jul 2024\)](#)

## 19.7 W3C Trace Context & Distributed Tracing

### Official W3C Specifications:

- [W3C Trace Context Specification](#)
- [W3C Trace Context Level 2](#)
- [GitHub: W3C Trace Context Repository](#)

### Implementation Guides:

- [OpenTelemetry: Propagation Guide](#)
- [Google Cloud: Trace Context Documentation](#)

## 19.8 Sensitive Data Redaction & PII Protection

### GDPR & Compliance:

- [Skyflow: How to Keep Sensitive Data Out of Your Logs \(Feb 2025\)](#)
- [Datadog: Redact Sensitive Data from Logs On-Prem](#)

### Industry-Specific:

- [OWASP Logging Cheat Sheet](#)
- [GDPR Official Text](#)
- [HIPAA Security Rule](#)

## 19.9 Uptime Kuma Integration

### Uptime Kuma Resources:

- [GitHub: Uptime Kuma Official Repository](#)
- [GitHub Wiki: Uptime Kuma Prometheus Integration](#)
- [Uptime Kuma Documentation](#)

## 19.10 Additional Resources

### OpenTelemetry Community:

- [OpenTelemetry Main Site](#)
- [OpenTelemetry Demo Applications](#)
- [CNCF OpenTelemetry Project](#)

### SigNoz Community:

- [SigNoz Slack Community](#)
- [SigNoz GitHub Discussions](#)
- [SigNoz Blog](#)

### ClickHouse Resources (For Reference):

- [ClickHouse Official Documentation](#)
- [ClickHouse Best Practices](#)

## 20. Glossary

| Term            | Definition  |
|-----------------|---|
| OTLP            | OpenTelemetry Protocol - Standard protocol for telemetry data |
| Span            | A single unit of work in a distributed system                 |
| Trace           | End-to-end journey of a request across multiple services      |
| Golden Signals  | Four key metrics: Latency, Traffic, Errors, Saturation        |
| Core Web Vitals | Google's metrics for page performance (LCP, FID, CLS)         |

|                             |   |
|-----------------------------|---|
| <b>MTTD</b>                 | Mean Time to Detection - Time to detect an issue            |
| <b>MTTR</b>                 | Mean Time to Resolution - Time to resolve an issue          |
| <b>RUM</b>                  | Real User Monitoring - Frontend performance tracking        |
| <b>PII</b>                  | Personally Identifiable Information                         |
| <b>Cardinality</b>          | Number of unique values for a given attribute               |
| <b>ClickHouse</b>           | Column-oriented database optimized for analytics            |
| <b>W3C Trace Context</b>    | Standard for propagating trace information across services  |
| <b>Semantic Conventions</b> | Agreed-upon naming standards for telemetry attributes       |
| <b>APM</b>                  | Application Performance Monitoring                          |
| <b>SLO</b>                  | Service Level Objective - Target for service performance    |
| <b>SLA</b>                  | Service Level Agreement - Contractual performance guarantee |
| <b>RBAC</b>                 | Role-Based Access Control                                   |
| <b>BAA</b>                  | Business Associate Agreement (for HIPAA compliance)         |

## 21. Quick Reference Guide

### Common Commands

#### OpenTelemetry Collector:

```
# Start collector
docker run -d --name otel-collector \
-p 4317:4317 -p 4318:4318 \
otel/opentelemetry-collector-contrib:0.93.0

# View collector logs
docker logs -f otel-collector

# Restart collector
docker restart otel-collector
```

### **SigNoz:**

```
# Access SigNoz Cloud  
https://your-org.signoz.cloud  
  
# API endpoint for data ingestion  
https://ingest.{region}.signoz.cloud:443
```

### **Uptime Kuma:**

```
# Start Uptime Kuma  
docker run -d --restart=always -p 3001:3001 \  
-v uptime-kuma:/app/data \  
--name uptime-kuma louislam/uptime-kuma:2  
  
# Access Uptime Kuma  
http://localhost:3001
```

## **Troubleshooting Checklist**

### **Logs not appearing in SigNoz:**

- [ ] Verify OTLP collector endpoint is correct
- [ ] Check collector is running (`docker ps`)
- [ ] Verify network connectivity to SigNoz Cloud
- [ ] Check API key/authentication is valid
- [ ] Review collector logs for errors
- [ ] Confirm log format is JSON structured

### **Traces not correlated:**

- [ ] Verify W3C Trace Context headers are present
- [ ] Check traceparent header format
- [ ] Ensure all services propagate trace context
- [ ] Verify trace ID is consistent across services
- [ ] Check OpenTelemetry SDK is properly initialized

### **High costs:**

- [ ] Review sampling rates (increase sampling for success logs)
- [ ] Check for duplicate log entries
- [ ] Verify log level is appropriate (not DEBUG in prod)
- [ ] Review metric cardinality
- [ ] Consider archiving old data to blob storage

#### **Slow queries in SigNoz:**

- [ ] Add indexes to frequently queried fields
- [ ] Reduce query time range
- [ ] Use filters to reduce data scanned
- [ ] Check if retention period is too long
- [ ] Contact SigNoz support for optimization

## **Emergency Contacts**

#### **On-Call Rotation:**

- Primary: [Your Team's On-Call Schedule]
- Backup: [Backup Schedule]
- Escalation: [Manager Contact]

#### **Vendor Support:**

- SigNoz Support: [support@signoz.io](mailto:support@signoz.io)
- SigNoz Slack: <https://signoz.io/slack>
- Emergency Hotline: [If applicable for Enterprise tier]

**Last Updated:** November 25, 2025